

```
In [39]: import numpy as np
import pandas as pd
import math as mth
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, accuracy_score, mean_absolute_error
```

```
In [2]: #made by yours truly, Mohamed Ehab
```

```
In [3]: df = pd.read_csv('advertising.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

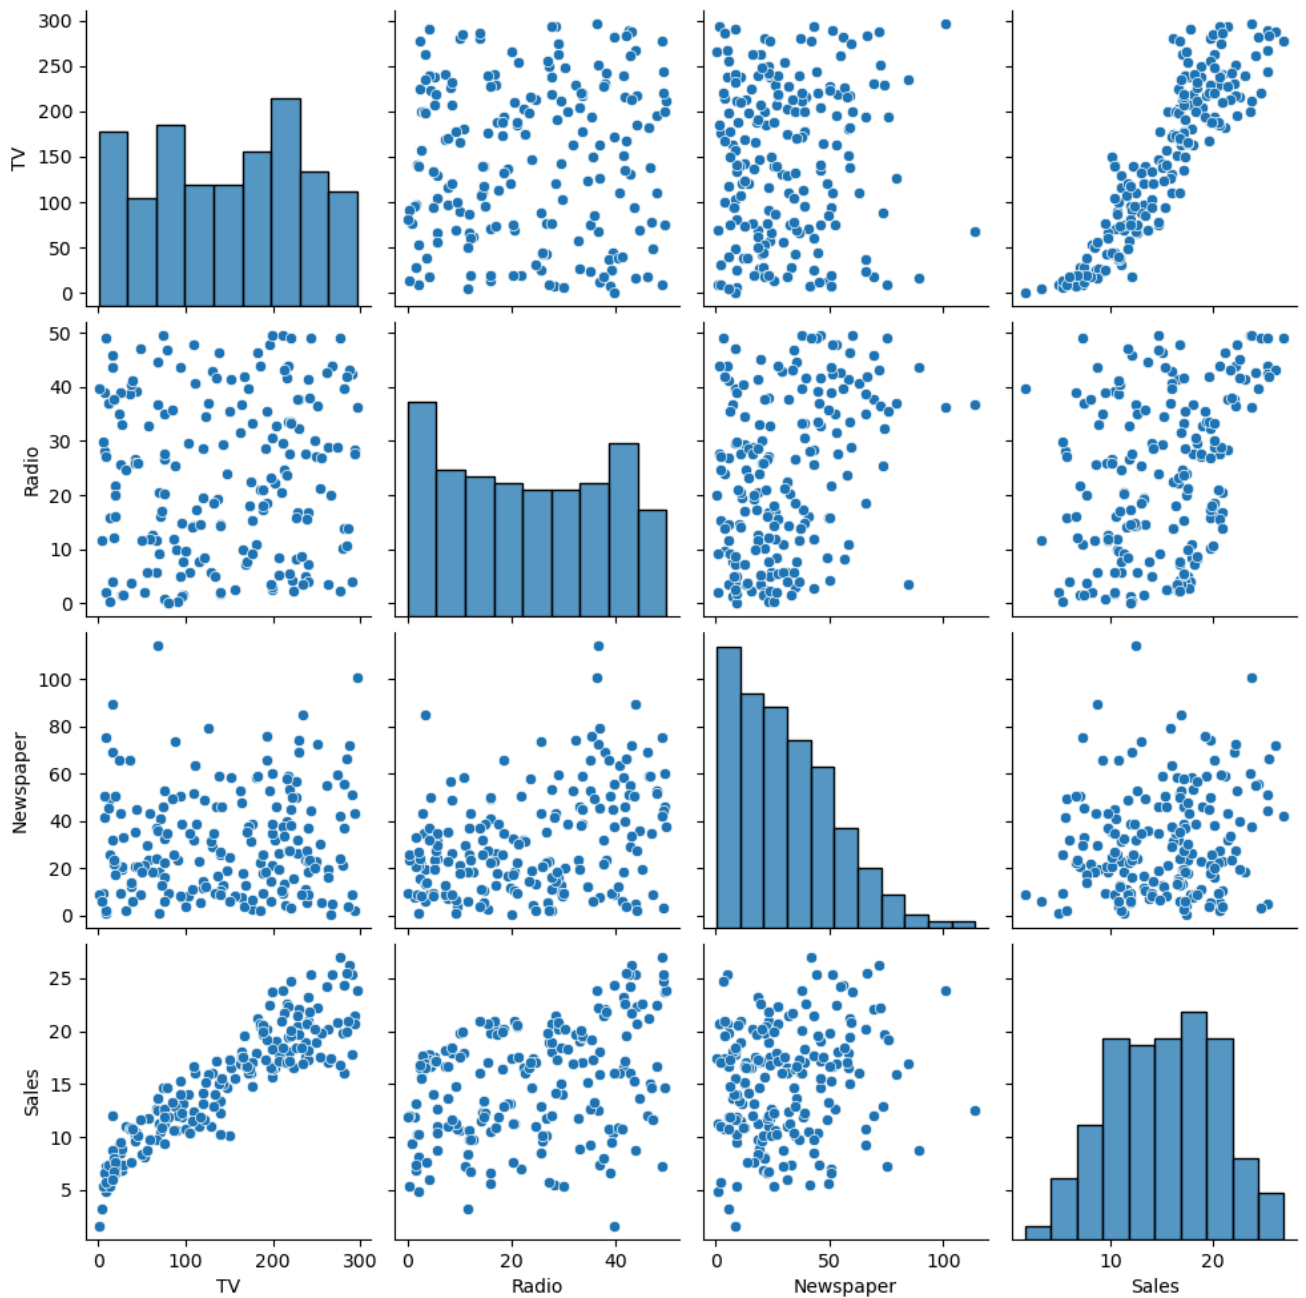
```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          196 non-null    float64
1    Radio       195 non-null    float64
2    Newspaper   195 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [38]: sns.pairplot(df)
```


C:\ProgramData\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)

```
Out[38]: <seaborn.axisgrid.PairGrid at 0x22a051e66d0>
```



```
In [37]: sns.kdeplot(df[['TV', 'Radio', 'Newspaper']])
```

```
Out[37]: <Axes: ylabel='Density'>
```



```
In [8]: print(f'How many duplicated values?: {df.duplicated().sum()}')
print(f'Null or missing values:\n{df.isnull().sum()}')
```

```
How many duplicated values?: 0
Null or missing values:
TV          4
Radio       5
Newspaper   5
Sales       0
dtype: int64
```

```
In [9]: df_clean = df.fillna(method='ffill')
```

```
In [10]: print(f'How many duplicated values after cleaning?: {df_clean.duplicated().sum()}')
print(f'Null or missing values after cleaning:\n{df_clean.isnull().sum()}')
```


```
How many duplicated values after cleaning?: 0
Null or missing values after cleaning:
TV          0
Radio       0
Newspaper   0
Sales       0
dtype: int64
```

```
In [11]: X = df_clean[[col for col in df.columns if col != "Sales"]]
y = df_clean[["Sales"]]
```

```
In [12]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

In [13]: `sns.kdeplot(X_scaled)`

Out[13]: `<Axes: ylabel='Density'>`



In [14]: `X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3)`  
`model = LinearRegression()`

In [15]: `model.fit(X_train, y_train)`

Out[15]: `LinearRegression`  
`LinearRegression()`

In [16]: `y_pred = model.predict(X_test)`


In [17]: `y_pred_int = y_pred.astype(int)`  
`y_test_int = y_test.astype(int)`

In [40]: `mse = round(mean_squared_error(y_test, y_pred),3)`  
`rmse = round(mth.sqrt(mse),2)`  
`accuracy = accuracy_score(y_pred_int, y_test_int)`  
`mae = round(mean_absolute_error(y_test, y_pred),3)`  
`print(f"Mean squared error: {mse}\nRoot mean squared error: {rmse}\nAccuracy score: {round(accuracy,3)}\nMean absolute error: {mae}")`

Mean squared error: 4.988  
 Root mean squared error: 2.23  
 Accuracy score: 0.3  
 Mean absolute error: 1.396

```
In [21]: residuals = y_test - y_pred

plt.scatter(y_pred, residuals)
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Sales')
plt.ylabel('Residuals')
plt.title('Residuals Plot')
plt.show()
```



although our errors are high, the scaled data we entered is a good fit for our regression model as we can conclude from **Residual plot**: the residuals are randomly scattered showing a linearity between the three features we entered.

our model can be far enhanced by **removing outliers** and **entering more data** for more accurate predictions