

```
In [95]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [94]: #made by yours truly, Mohamed Ehab
```

```
In [2]: df = pd.read_csv("Bank Customer Churn Prediction.csv")
```

```
In [4]: df.head()
```

Out[4]:

|   | customer_id | credit_score | country | gender | age | tenure | balance   | products_number | credit_card |
|---|-------------|--------------|---------|--------|-----|--------|-----------|-----------------|-------------|
| 0 | 15634602    | 619          | France  | Female | 42  | 2      | 0.00      | 1               | 1           |
| 1 | 15647311    | 608          | Spain   | Female | 41  | 1      | 83807.86  | 1               | 0           |
| 2 | 15619304    | 502          | France  | Female | 42  | 8      | 159660.80 | 3               | 1           |
| 3 | 15701354    | 699          | France  | Female | 39  | 1      | 0.00      | 2               | 0           |
| 4 | 15737888    | 850          | Spain   | Female | 43  | 2      | 125510.82 | 1               | 1           |

```
In [6]: df = df.drop(columns=["customer_id"])
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   credit_score           10000 non-null   int64  
1   country                10000 non-null   object  
2   gender                 10000 non-null   object  
3   age                   10000 non-null   int64  
4   tenure                 10000 non-null   int64  
5   balance                10000 non-null   float64 
6   products_number        10000 non-null   int64  
7   credit_card            10000 non-null   int64  
8   active_member          10000 non-null   int64  
9   estimated_salary        10000 non-null   float64 
10  churn                  10000 non-null   int64  
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

```
In [15]: print(df.isnull().sum())
print("\nduplicated: ",df.duplicated().sum())
```

```
credit_score      0
country           0
gender            0
age              0
tenure           0
balance           0
products_number   0
credit_card       0
active_member     0
estimated_salary  0
churn             0
dtype: int64
```

```
duplicated: 0
```

```
In [20]: df['country'].unique()
```

```
Out[20]: array(['France', 'Spain', 'Germany'], dtype=object)
```

```
In [21]: df['products_number'].unique()
```

```
Out[21]: array([1, 3, 2, 4], dtype=int64)
```

```
In [26]: df = pd.get_dummies(df, drop_first=True)
```

```
In [41]: df.head()
```

```
Out[41]:
```


|   | credit_score | age | tenure | balance   | products_number | credit_card | active_member | estimated_sala |
|---|--------------|-----|--------|-----------|-----------------|-------------|---------------|----------------|
| 0 | 619          | 42  | 2      | 0.00      | 1               | 1           | 1             | 101348.        |
| 1 | 608          | 41  | 1      | 83807.86  | 1               | 0           | 1             | 112542.        |
| 2 | 502          | 42  | 8      | 159660.80 | 3               | 1           | 0             | 113931.        |
| 3 | 699          | 39  | 1      | 0.00      | 2               | 0           | 0             | 93826.         |
| 4 | 850          | 43  | 2      | 125510.82 | 1               | 1           | 1             | 79084.         |

```
In [42]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   credit_score          10000 non-null  int64  
 1   age                   10000 non-null  int64  
 2   tenure                10000 non-null  int64  
 3   balance               10000 non-null  float64 
 4   products_number       10000 non-null  int64  
 5   credit_card           10000 non-null  int64  
 6   active_member         10000 non-null  int64  
 7   estimated_salary      10000 non-null  float64 
 8   churn                 10000 non-null  int64  
 9   country_Germany       10000 non-null  bool    
10   country_Spain         10000 non-null  bool    
11   gender_Male           10000 non-null  bool    
dtypes: bool(3), float64(2), int64(7)
memory usage: 732.6 KB
```

```
In [32]: df.churn.plot.hist()
```

```
Out[32]: <Axes: ylabel='Frequency'>
```



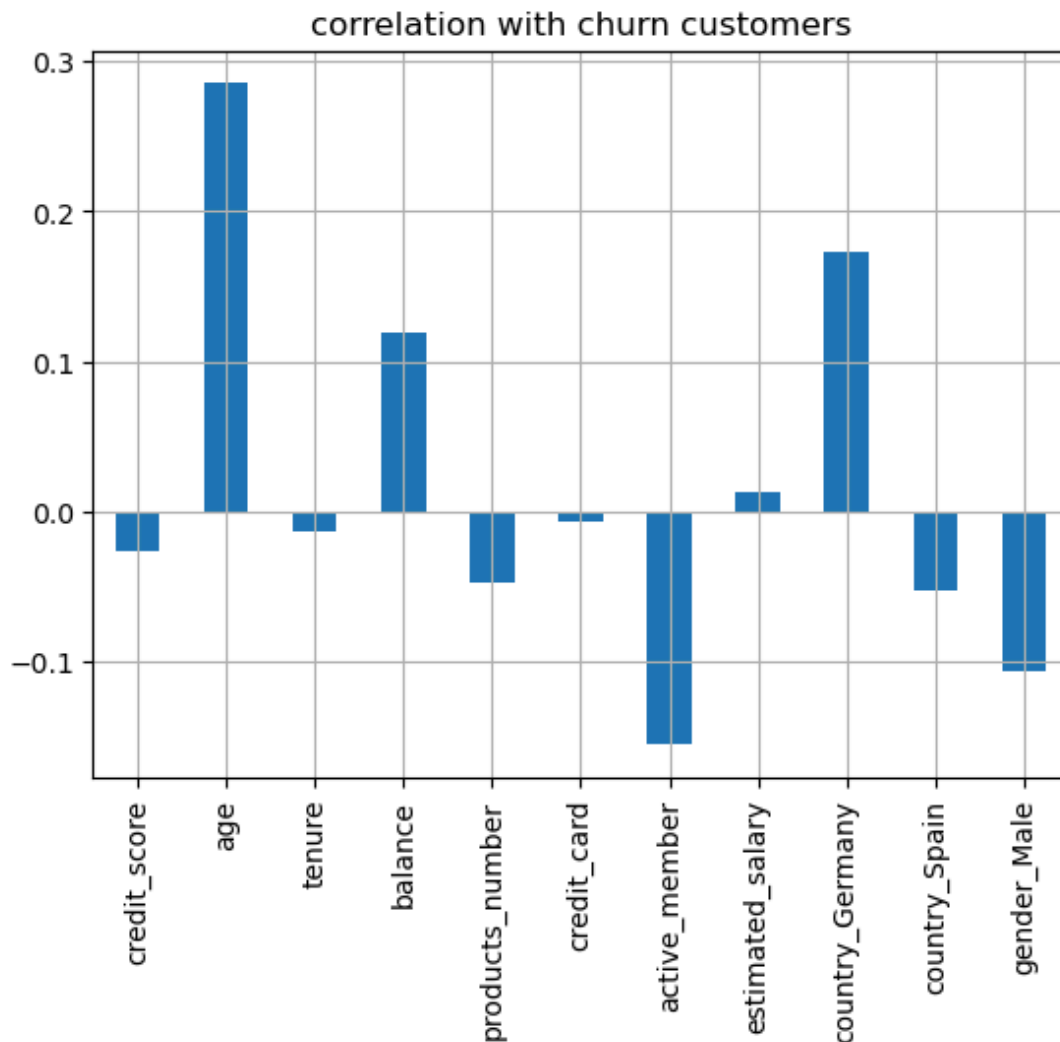
```
In [37]: loyal = (df.churn==0).sum()
churned = (df.churn==1).sum()
print(f"{loyal} is still bank customers\n{churned} has left the bank.")
```

```
7963 is still bank customers
2037 has left the bank.
```

```
In [56]: X = df.drop(columns='churn')
y = df['churn']
```


```
In [57]: X.corrwith(df['churn']).plot.bar(title='correlation with churn customers', grid=True)
```

```
Out[57]: <Axes: title={'center': 'correlation with churn customers'}>
```



```
In [51]: sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt='.2f')
```

```
Out[51]: <Axes: >
```



```
In [68]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [71]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [72]: model = LogisticRegression()
model.fit(X_train, y_train)
```


```
Out[72]: LogisticRegression
LogisticRegression()
```

```
In [73]: y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {:.2f}%".format(accuracy * 100))
```

Accuracy: 80.15%

```
In [93]: sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='.0f')
```

Out[93]: <Axes: >



```
In [83]: print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.82         0.96     0.88         1581
     1       0.58         0.20     0.30          419

 accuracy          0.80
 macro avg         0.70         0.58     0.59
 weighted avg      0.77         0.80     0.76
```

In [ ]:

